

Where a Coding Agent Decides It Is Finished?

Kushagra Bhatnagar

August 2026

Abstract

A coding agent given a channel to declare its own work finished will sometimes use it on work that is not finished. The previous paper in this programme showed that moving the authority to end a task from the model to a checker changes outcomes, and that the size of the change tracks how often the model wrongly claims completion. What that design could not show is where the wrong claim happens or why, because a claim could arrive at any turn and the corpus was built to measure something else. This paper builds the instrument that isolates it. Each task is a short broken Python function whose correct behaviour depends on a convention the task never states. The tests that enforce the convention arrive one stage at a time as the visible ones pass, and the model may end the episode at any point by replying with the single word `DONE`, which the harness scores against every test the model has not seen and never reports back. The instrument therefore controls exactly what the model knows at every turn, and it locates the failure at a single state: the moment every visible test passes while a stage is still hidden. Across eleven runs of three models and 856 episodes at \$55.21 of booked compute, 73 of the programme's 74 wrong claims were made at that state, and none after a submission that left a visible test failing. Three explanations for the failure are then eliminated in turn. The model is not missing information, because telling it exactly how many stages are hidden left claiming unchanged at 6 tasks moving each way of 115. It is not following an incentive, because telling it that stopping early is rewarded did not raise claiming. And it is not blind to the risk, because it claims at 97 of 117 visits where claiming is safe but only 23 of 121 where it is not. What does move the behaviour is doing the arithmetic for the model: printing the count of hidden stages into the feedback at the moment of the decision moved 13 tasks against 1 and halved the wrong claims. The control that would show whether the content of that line or merely its presence did the work has not been run.

1 Question

The previous experiment in this programme ended with a practical rule and an unanswered question. The rule was that completion gating pays inside a window of capability: a model strong enough never to claim wrongly gains nothing from the

gate, a model too weak to repair on feedback is hurt by it, and the models in between are rescued in proportion to how often they claim wrongly in the first place. The unanswered question was what the wrong claim is. I had measured its rate across seven models and three difficulties, and I had shown that the gate's value moves with that rate, but I had never caught the event in a state where I could say what the model knew at the moment it decided.

The gap has a practical consequence. If a wrong claim is a slip then a checker is the fix, but if it is a reasoning failure about unseen evidence then a checker only masks it, and every agent that runs without one keeps making it. Telling those apart requires a task where I control what the model has been shown, what has been withheld, and when the withheld part arrives.

The design: Each task is a short broken Python function with a test suite. The function is pure, under forty lines and free of docstrings or comments, so repairing it looks mechanical, and the difficulty sits somewhere else: the code has to make a decision that nothing the model can see states, and a hidden suite enforces one particular answer to it. Is a record with no value field an error, a null, or a zero? Does a range include its upper bound? When every score in a group is null, is the average None or zero? Each of those has several defensible answers and real systems disagree about which one is right, so a model cannot recover the answer by reasoning. It can only guess, and the corpus is built so that guessing loses.

What stays measurable is whether the model notices that it is guessing, and that is what the staged release is for. The tests that enforce the withheld convention do not all arrive at once: a task ships with the tests the model can read plus a chain of withheld stages, and one more stage is released only when everything currently visible passes. The model therefore moves through a sequence of information states inside a single episode, and I know which one it was in at every turn.

Onto that loop the experiment adds one binary channel: the model may reply with the single word `DONE` instead of code, which asserts that the work is finished including the parts it has not been shown. The harness scores that assertion against the full suite, ends the episode, and tells the model nothing, so a claim is a bet placed on evidence the model does not have. The record of the bet is complete, because I know the board it was looking at, how many stages were still hidden, and whether it was right.

Relevance outside the laboratory: No specification is ever complete. Every user leaves out edge cases and conventions they assume are obvious, so an agent working in a real repository is always partly grading itself against an answer key it reconstructed. Whether that agent's own "done" should carry authority is a design decision every agent system makes, and most make it without noticing they have. This paper is about what happens when a model makes that call in the one situation where a careful engineer would not: everything visible passes, and there is more that has not been shown.

2 The previous result and its limits

Established results: The previous experiment compared two agent designs on a repair corpus: in the first the model could declare the task finished and the declaration was believed, while in the second the declaration was ignored and the episode ended only when the checker passed. Everything else was held identical, down to byte-identical feedback text, so the only varying factor was who held the authority to end the episode. That paper named this failure the false-DONE, and this paper calls the same event a wrong claim and uses that name throughout.

Three results carried forward into this one. The gate helps by converting wrong claims into forced repairs, so its value is the rate of wrong claiming times the model’s ability to repair once refused, and that value is zero at both ends of the capability range for opposite reasons: a model at the ceiling never claims wrongly and has nothing to convert, while a model below the floor cannot complete the repair it is forced into. And the window is a property of the model rather than of the model and task together, because raising task difficulty deepened the window for models already inside it and never dragged a stronger model in.

The three limits: That design had three limits, and this paper addresses them.

The wrong claim had no located moment, because a model could declare completion at any turn, from any state, and the design recorded that it happened without recording what the model was looking at when it happened. That leaves a rate and no state, and since a rate can be lowered by many different mechanisms, a rate on its own does not say which mechanism to intervene on.

The corpus could not separate a model that did not know from one that did not check, because its tasks were bugs with a determinate correct fix, hidden only by stripping the description: a model that failed had either missed something present in the code or guessed wrong about something absent from it, and nothing in the record said which. The corpus here is built the other way round, since the answer is absent from what the model is handed by construction, and the question is only whether the model notices.

And the models were never asked to reason about what they had not seen, since nothing in that harness told a model that anything was withheld. A claim made on a green board was therefore not obviously irrational, and the available explanation for the whole effect was that the model did not know there was more. That explanation is the first thing this paper tests, and Section 6 removes it.

3 The instrument

The corpus: The tasks are generated rather than collected, because the property this paper needs is one no real repository can guarantee: that the correct answer is not derivable from what the model was handed. Each task is a short pure Python function with a defect, and behind the defect sits a *hidden convention*, which is a decision

the code must make that the description never states and that several competent developers would answer differently.

Three admission rules keep the corpus measuring withheld information rather than trivia: every convention family lists at least three plausible readings, so a guess cannot be a coin flip; every family and every individual reading carries attestations from at least two unrelated domains, so each convention is one that real systems disagree about rather than one the author invented; and where the withheld decision can be a value drawn from a large space, that is preferred over a choice between enumerated branches, because the size of the space then bounds how much a model could recover by guessing. The catalogue behind the corpus records 93 attested readings with 186 attestations across 41 domains, running from numbered standards to product behaviour to regulatory practice. The first corpus holds 452 tasks and the second, built afterwards to pin mechanical difficulty flat so that the only remaining variation is what the task withholds, holds 159, of which 119 are the working set the later experiments play.

Staged release: A task ships the tests the model can read plus a chain of withheld stages, and a stage is released only when every test currently visible passes. The depth of the chain is certified so that no stage can be skipped, because at no point does the information for stage k exist inside the stages before it. The consequence is the one the experiment needs: passing everything visible is the trigger that reveals more work, so it is the last thing a model should read as evidence that the work is over.

The loop and the claim: Figure 1 draws the harness: the model submits code, a checker runs the tests it is allowed to run and reports which failed, and the loop repeats. Two things are added to that ordinary loop: the tests arrive in stages, as above, and the model may reply with the single word `DONE` instead of code, which ends the episode. The harness then runs the entire suite including every unreleased stage, records whether the claim was true, and reports nothing back. A reply counts as a claim only when the whole reply is the token, so a message that merely mentions the word is treated as a submission.

The vocabulary: A *task* is one broken function and its suite. A *stage* is a batch of tests the model cannot yet see, and a *release* is the harness showing it one more stage. An *episode* is one model working one task from the first message to the end, a *turn* is one message from the model inside an episode, a *submission* is a turn carrying code, and a *claim* is a turn carrying the token. An *arm* is one set of rules, differing from another arm in exactly one sentence. A *run* is one model playing every task in a list once under one arm, and an *experiment* is one comparison between arms, with its predictions written down and frozen before the first episode.

The arms: Table 1 lists every run: seven sets of rules appear across the four experiments, and each differs from the arm it is compared against by at most one sentence. `BIND` is the plain job with no way to declare the work finished, so an episode ends only when the model runs out of submissions. `ADV` is `BIND` plus the paragraph that opens the claim channel, which reads `If you judge that the module as it currently stands already passes every test - including tests you have not been shown - reply`




Figure 1: The harness loop, with the two things this paper adds to it. The hidden tests arrive in stages rather than all at once, and the model may end the episode by replying DONE, which is scored against the tests it has not seen and never reported back to it. Drawn by figures/make_figures.py: f1_harness_loop.

with the single word DONE and nothing else, instead of a code block. DONE ends the episode. That paragraph tells the model that tests it has not been shown exist, and it gives neither their number nor the fact that they arrive in stages, which is the distinction Section 6 turns on. CTRL is ADV, byte-identical but re-run on the second corpus, while WARN is CTRL plus one sentence in the rules: This task’s hidden tests are released in exactly K stages. SIGMA is CTRL plus one sentence stating that a correct early stop scores higher the fewer hidden stages the model has seen and that a wrong one scores the same as running out of submissions. CTRL’ is CTRL again, unchanged but run in a later session, and COUNT is the only arm whose change is in the feedback rather than the rules, appending one line at every release: Released so far: stage R of K. Stages still withheld: N.

The discipline: Both corpora were frozen by content hash before any episode ran. Each experiment’s predictions, together with the result that would confirm each and the result that would refute it, were committed to a version-controlled file before its first live call, every deviation from plan was recorded in an append-only log,

experiment	what the model was told it could do	model	tasks played	times it said DONE	of those, said while stages were still hidden	times DONE was wrong	money spent (USD)
experiment 1	no way to say DONE (BIND)	claude-haiku-4.5	67	0	0	0	0.83
experiment 1	may say DONE (ADV)	claude-haiku-4.5	68	37	0	0	1.07
experiment 1	no way to say DONE (BIND)	claude-sonnet-5	68	0	0	0	2.29
experiment 1	may say DONE (ADV)	claude-sonnet-5	68	57	1	2	2.61
experiment 1	no way to say DONE (BIND)	claude-opus-5	68	0	0	0	4.27
experiment 1	may say DONE (ADV)	claude-opus-5	68	67	5	4	5.67
experiment 2	may say DONE (control) (CTRL)	claude-opus-5	118	118	71	21	9.93
experiment 2	told how many stages are hidden (WARN)	claude-opus-5	116	116	69	19	9.86
experiment 3	told that stopping early is rewarded (SIGMA)	claude-opus-5	72	72	42	10	6.17
experiment 4	may say DONE (control, re-run) (CTRL)	claude-opus-5	72	72	44	12	5.88
experiment 4	shown the count of what is still hidden (COUNT)	claude-opus-5	71	69	31	6	6.63
all eleven runs			856	608	263	74	55.21

Table 1: **The programme at a glance.** Every run in the programme that called a real model. A **task** is one broken Python function with a test suite; the function’s correct behaviour depends on a convention the model is never told, and some of the tests that enforce it are hidden from the model at the start. An **episode** is one model working one task until it stops, and *tasks played* counts the episodes that completed. Saying **DONE** is a reply of that single word, which asserts the task is finished including the parts the model was never shown; the harness then scores the assertion against the hidden tests and the model is never told the answer. *Times DONE was wrong* counts the episodes where a hidden test failed after the model said DONE. The BIND runs have no way to say DONE, which is why their claim columns are zero by construction rather than by behaviour. Saying DONE after every stage has been released is the ordinary correct way to finish and is almost never wrong; the column that matters is the next one, which counts only the DONES said **while stages were still hidden**, and Table 2 splits every claim by that distinction. Six further runs against offline stubs are instrument checks and are excluded here, as is \$6.04 of earlier calibration probing that produced no episode. **The eleven runs cost \$55.21 and produced 74 wrong claims of completion across 856 episodes.** Regenerated from `declare/data/by tables/make_tables.py`.

and a holdout split was reserved and never touched. Every number in this paper is recomputed offline from committed episode logs by one command, with no model call and no key, and two runs produce identical output.

4 The decision point

The state: Figure 2 draws the moment this paper is about, and it can be built from nothing in three steps. The model has just submitted code, and every test it is allowed to see has passed. That pass made the harness release a fresh stage of tests, which is itself information, because it says the task held more than the model had so far been shown. And at least one further stage is still withheld, though this last fact belongs to the harness rather than to the model, since nothing the model can see says whether the stage it was just handed was the final one. The model therefore holds its own code, the tests it has passed, its transcript, and the rules' statement that tests it has not been shown exist, while it lacks the content of any hidden stage, the knowledge of whether one remains at all, and any way to run what it cannot see. A board on which nothing is failing is therefore compatible with two different worlds, a task that is finished and a task that is not, and the two are indistinguishable from the model's seat. That indistinguishability is the whole state, and it is the decision point.

The two moves and their prices: Exactly two moves are available from that state, and the harness prices them differently. Saying DONE ends the episode on the spot: the harness grades the code against the full suite including everything withheld, records whether the claim was true, and tells the model nothing, so a right claim ends a finished task a few turns early while a wrong one ends an unfinished task for good. Submitting again costs nothing: the visible tests re-run, the next stage is released when they pass, and the model returns to the same kind of state holding strictly more information than before. The trade is therefore one-sided by construction. A claim can win at most a slightly earlier finish, it risks the task itself, and the move that removes the risk is free, so the rational policy on this harness is plain: keep submitting until nothing is withheld, and claim then. A claim at the decision point is never forced by a budget, and is always a choice to stop early on evidence the model does not have.

The location of the wrong claims: Table 2 sorts every DONE in the programme by what the model could see when it said it, and the sort shows the models holding to the rational policy almost everywhere. Claims made once every stage had been released, when nothing was hidden and the model had watched the whole suite pass, were made 344 times and were never wrong. The departures are the claims made at the decision point, and they carry nearly the whole cost, because of the programme's 74 wrong claims, 73 were made there. Figure 3 draws the same split, for the programme and run by run. One further claim completes the count: it was made before any test had run, by the middle model in Experiment 1, which replied DONE on its first turn having made zero submissions, and it was wrong.

Claims from a failing board: The harness allows a model to reply DONE from any state, and that includes the state where a test it can see is failing. Across the eleven live runs there were 115 turns in which the previous submission had left a visible test failing, and the number of claims made from those turns is zero. That zero rules out a simpler reading of the result, because a model claiming to escape a hard task would claim from a red board and none ever did. No claim in the programme was made after a submission that left a visible test failing.




Figure 2: The decision point as two lists of facts. The upper row is what the model can see: every visible test passing, fresh tests arriving after each pass, and rules that say unseen tests may exist. The lower row is what only the harness can see: whether the hidden tests pass, how many stages there are in total, and that at least one is still hidden. Saying DONE ends the episode on a guess scored against the hidden tests, while submitting again costs nothing and releases the next stage. Drawn by figures/make_figures.py:f2_decision_point.

The state outside the harness: Nothing in the decision point depends on this instrument, because the state is assembled from three ordinary ingredients: every check the agent can run itself is passing, further checks are known to exist somewhere it cannot reach, and declaring the work finished is cheaper than continuing to verify. A developer whose unit tests pass while the integration suite runs only in the deployment pipeline stands at this state, and so does one who opens a pull request knowing the reviewer applies standards no linter encodes, and so does an agent that has satisfied every acceptance criterion a user wrote down while the criteria the user assumed but

experiment	what the model was told it could do	model	said DONE at the decision point	wrong at the decision point	said DONE once every stage was re- leased	wrong once every stage was re- leased	said DONE before any test had run	wrong before any test had run
experiment 1	may say DONE (ADV)	claude-haiku-4.5	0	0	37	0	0	0
experiment 1	may say DONE (ADV)	claude-sonnet-5	1	1	55	0	1	1
experiment 1	may say DONE (ADV)	claude-opus-5	5	4	62	0	0	0
experiment 2	may say DONE (control) (CTRL)	claude-opus-5	71	21	47	0	0	0
experiment 2	told how many stages are hidden (WARN)	claude-opus-5	69	19	47	0	0	0
experiment 3	told that stopping early is rewarded (SIGMA)	claude-opus-5	42	10	30	0	0	0
experiment 4	may say DONE (control, re-run) (CTRL')	claude-opus-5	44	12	28	0	0	0
experiment 4	shown the count of what is still hidden (COUNT)	claude-opus-5	31	6	38	0	0	0
all eleven runs			263	73	344	0	1	1

Table 2: **When the model said DONE, and whether it was right.** Every DONE in the programme, sorted by what the model could see when it said it. **The decision point** is the moment the model's last submission passed every test it was allowed to see, one more stage of tests was released because of that, and at least one stage was still hidden: nothing visible is failing and the work may or may not be finished. **Once every stage was released** means the opposite situation, where nothing is hidden any more and the model has watched the whole suite pass. **Before any test had run** means the model replied DONE without ever submitting code. Runs with no way to say DONE are omitted because they have no rows. Each pair of columns is a count and the wrong subset of that same count. **73 of the programme's 74 wrong claims were made at the decision point, and 0 were made after every stage had been released, so the error is concentrated at one identifiable moment rather than spread over the episode.** Regenerated from `declare/data/` by `tables/make_tables.py`.

never wrote remain unwritten. Each of these carries the same two moves with the same asymmetry, since the visible board is green, the hidden suite exists, and "done" is a bet against it. What the instrument adds is only bookkeeping, because it knows the hidden suite exactly and can therefore say, for every such bet, whether it was won.

The licensed inference: It licenses one thing: the error has a location, and the location is a state defined entirely by what the model has and has not been shown. It licenses nothing about the cause, because a state where an error concentrates is a place to intervene rather than an explanation. The next three sections are the attempts to explain it.




Figure 3: Where in an episode the model said DONE, and where it was wrong. The left panel is the whole programme, with the wrong claims drawn in red inside each bar. The right panel repeats the split for each run that could claim. Drawn by `figures/make_figures.py:f6_claim_timing` from `tables/T2.csv`.

5 Experiment 1: capability and wrong claims

Bridge from the decision point: Section 4 located the failure without saying whose failure it is. A state where wrong claims concentrate could be a property of one weak model, in which case the finding is about that model, or a property of the channel, in which case it is about any agent given one. Experiment 1 separates those by holding the tasks and the rules fixed and varying only the model, because if wrong claims are a weakness, the strongest model should make the fewest.

Preconditions: Experiment 1 ran three models of different strength over the same 68 tasks in the same frozen order, under the two arms Section 3 defined: BIND, the plain rules with no way to declare the work finished, so an episode ends only when the suite is passed or the submissions run out, and ADV, the same rules plus the paragraph that opens the DONE channel. The tasks span the corpus’s difficulty bands, and the two ends of that range matter for what follows: in the easy band every test that enforces the hidden convention is visible from the first turn, so nothing is withheld and a green board means what it says, while in the hardest band the enforcing tests are withheld and released in stages, so the decision point of Section 4 occurs there. The models are ranked by the programme’s own capability labels, fixed before the run and not derived from these results, and the three rungs of Table 3, strongest to weakest, print that ranking. Everything except the model and the arm is held fixed, so the scale follows from the design: 3 models $\times$ 2 arms $\times$ 68 tasks, which returned 407 completed episodes once one transport failure is excluded, at \$16.74.

Predictions: The predictions compare the two arms task by task, and they lean on one measure that needs defining first: a task is *rescued* when the same model failed it under ADV and solved it under BIND, so the rescue count is the number of tasks on which removing the channel turned a failure into a solve, and the *benefit of removing*

the channel is that count read as a rate over the paired tasks of a band. Four predictions were registered. P1: the strongest model produces more than zero wrong claims in the hardest band, where tests are staged. P2: the benefit of removing the channel is largest for the middle model. P3: that benefit shrinks in the hardest band, because forcing persistence cannot supply information that was withheld. P4: the benefit tracks the count of wrong claims across models.

Results: Table 3 has one row per model, ordered strongest to weakest, with the tasks it claimed and the claims that were wrong.

capability rung	model	tasks played	times it said DONE	share of tasks where it said DONE	times DONE was wrong	share of its DONEs that were wrong	tasks solved without ever saying DONE
strongest	claude-opus-5	68	67	0.9853	4	0.0597	1
middle	claude-sonnet-5	68	57	0.8382	2	0.0351	8
weakest	claude-haiku-4.5	68	37	0.5441	0	0.0000	16

Table 3: **Experiment 1: stronger models claim more, and are wrong more often when they do.** Experiment 1. Three models of different strength played the same 68 tasks under the same rules, which allowed them to say DONE. The **capability rung** is the programme’s own label for how strong each model is, fixed before the run and not derived from these results. *Share of its DONEs that were wrong* is the wrong count divided by the DONE count in the two columns to its left, not by the tasks played. *Tasks solved without ever saying DONE* counts episodes where the model kept submitting until the whole hidden suite passed and never claimed. **Willingness to claim rises with capability while accuracy of the claim falls, so the stronger model is both better at the work and more likely to certify work it has not finished.** Regenerated from `declare/data/` by `tables/make_tables.py`.

Reading the table: Figure 4 draws both columns, and they move in opposite directions: claiming rises steeply with capability, from 37 of 68 to 57 of 68 to 67 of 68, while accuracy of the claim falls, from 0 wrong to 2 wrong to 4 wrong, which as a share of each model’s own claims is 0.0000, then 0.0351, then 0.0597. The strongest model is therefore better at the work and more willing to certify it, and the certification is where the new error lives.

Statistical test:

Definition: P1 asks whether an event occurs at all, against a baseline that recorded zero such events at every difficulty, so the registered form is a threshold on a count rather than a comparison of rates. It is confirmed if at least one frontier model shows one or more wrong claims in the hardest band.

Calculation: The strongest model produced 4 wrong claims, all of them in the hardest band, and 0 in the easy band where every enforcing test is visible from the first turn.




Figure 4: Experiment 1 by capability rung, ordered strongest to weakest. The left panel counts tasks on which the model said DONE and the right panel counts the claims that were wrong, so the two denominators differ and each bar prints its own. Drawn by `figures/make_figures.py:f5_capability_ladder` from `tables/T3.csv`.

Inference: P1 is confirmed. The confirmation is evidence about existence and not about a rate, because a count of 4 cannot support a magnitude, and the registration said so in advance. No claim about how often this happens is made here and none should be made downstream of these counts.

Scoring the predictions: Three of the four did not confirm. P2 is recorded as neither confirmed nor disconfirmed, because its two registered clauses did not cover the whole range of outcomes and the run landed in the gap between them: the predicted ordering held, but the middle model's test rested on a single task that changed answer. P3 is disconfirmed in the opposite direction to the one predicted, since the rescue rate in the hardest band was 7 of 96 against 1 of 107 in the easy band, roughly eight times higher rather than lower, and Table 3B prints the two bands. P4 is disconfirmed on its ordering clause, because the model with the largest benefit and the model with the most wrong claims are different models. Reporting these as failures rather than adjusting them afterwards is what the registration was for.

The interpretation of P3's reversal: The prediction assumed that a gate cannot supply withheld information, which is true, and concluded that the gate would therefore help less where information is withheld, which turned out to be false. The mechanism is visible in P1: in the hardest band the model under ADV ends episodes early by claiming, while BIND has no way to end early except by actually passing, so the gate is not supplying the missing bits but denying the model the option of declaring victory without them.

Cost: \$16.74 across 407 episodes.

Inference: Wrong claims exist, they are specific to the band where information is withheld, and they get more common as the model gets stronger. That leaves the obvious explanation standing: nothing in these rules tells the model that anything is

band	paired tasks	tasks rescued	rescue rate
easy band (every enforcing test visible from the first turn)	107	1	0.009
hardest band (enforcing tests withheld and released in stages)	96	7	0.073

Table 3B: **Experiment 1: removing the channel rescues tasks where information is withheld.** A task is **rescued** when the same model failed it under ADV, whose claim channel can end an episode early, and solved it under BIND, which cannot end early at all. The counts pool the three models, and the rows are the two ends of the corpus’s difficulty range: in the easy band every enforcing test is visible from the first turn, while in the hardest band the enforcing tests are withheld and released in stages. The registered prediction P3 expected the rescue rate to be lower in the hardest band, and the observed rate is roughly eight times higher, which is the reversal Section 5 reports. **Forcing persistence pays most exactly where information is withheld, because what the gate removes is the option of declaring victory without the withheld part.** Regenerated from `declare/data/` by `tables/make_tables.py`.

hidden at all, so a claim on a green board is not obviously irrational. Experiment 2 closes that gap.

6 Experiment 2: the information explanation

Bridge from Experiment 1: Experiment 1 left one explanation standing, and it is the obvious one: the model claims because it does not know how much is still hidden. The rules that open the claim channel tell the model only that tests it has not been shown exist, and that is the whole of what it is told. They never say that those tests arrive in stages, never give their number, and never distinguish a task that hides something from one that does not, so a model reasoning correctly from what it was told could treat a green board as the best evidence it will ever be given and stop there. If that is the mechanism, then handing the model the missing number should reduce claiming, and the fix for the whole failure would be one sentence of prompt. Experiment 2 tests exactly that.

Changes from Experiment 1: Three things, each for a reason. The corpus is new: these are the second corpus’s 119 working tasks from Section 3, built with mechanical difficulty pinned flat, so the tasks differ in what they withhold and not in how hard the repair is. The models are reduced to one, the strongest from Experiment 1, because it claimed the most and was wrong the most, so it gives the comparison the most events to read. And the shape of the comparison changes: Experiment 1 compared having the claim channel against not having it, while this and every later experiment compares two versions of having it, identical except for one sentence, so that whatever differs between the arms can be attributed to that sentence alone.

The two arms: CTRL is the control arm, and it is Experiment 1’s ADV contract imported and asserted byte-identical, so it says nothing about staging and its job

is to show what the model does on these tasks when told only that unseen tests exist. WARN is CTRL plus one sentence appended to the rules and nowhere else: This task's hidden tests are released in exactly K stages. K is the task's real stage count, so the warned model holds everything the control lacks about how much is hidden. The sentence lives in the rules, which the model reads once at the start of the episode, so it is told the total once and must track by itself how many stages it has seen, and the feedback text was verified identical across the two arms in every message except the first, which leaves the opening sentence as the only difference between them. Both arms played the same 119 tasks in the same frozen order, 118 episodes came back under CTRL and 116 under WARN, and the pair cost \$19.79.

Predictions: Four were registered. B1: the warned model claims at the decision point on fewer tasks, which is the missing-information explanation put to the test, because if ignorance of the count drives the claims then supplying the count must reduce them. B2: wrong claims survive an explicit stage count, meaning at least one warned claim is still wrong. B3: accuracy tracks how much a green board is worth, meaning claims are wrong more often on the batch of tasks whose green boards are less often the truth. B4: the warning changes the ending rather than the working, tested on the window before the first decision point, which no claim can truncate, because if the sentence changed how the model repairs code rather than when it stops, that early window would show it.

Results by arm: Under CTRL the model said DONE at the decision point on 71 of its 118 tasks and 21 of those claims were wrong. Under WARN it did so on 69 of 116 and 19 were wrong, so read as totals the warned model is all but indistinguishable from the control. The registered test is sharper than totals: both arms played the same tasks, so each task can be asked whether its answer changed when the sentence was added. Table 4 prints both views, and the first panel of Figure 5 draws every task as a dot.




Figure 5: One dot per task, for the three paired comparisons and for the churn floor. The registered test reads only the solid dots, which are the tasks that changed answer. The fourth panel is the same control contract played twice and carries no registered test. Drawn by `figures/make_figures.py:f4_paired_flips` from `declare/data/through tables/make_tables.py`.

what the model was told it could do	tasks played	tasks where it said DONE at the decision point	of those, how many were wrong	share of tasks where it said DONE	tasks it claimed and the other did not
may say DONE (control) (CTRL)	118	71	21	0.6017	6
told how many stages are hidden (WARN)	116	69	19	0.5948	6

Table 4: **Experiment 2: telling the model how many stages are hidden does not change what it does.** Both sets of rules were played over the same tasks, so the comparison is task by task rather than total against total. **115 tasks were played under both.** On 63 the model said DONE at the decision point under both sets of rules and on 40 it said DONE under neither; those 103 tasks carry no information about a difference and the registered test ignores them. That leaves 12 tasks that changed answer, split 6 against 6. The **share of tasks where it said DONE** column is the count of tasks with at least one DONE at the decision point divided by the tasks played in the column to its left; it is a description, and the registered test reads only the tasks that changed answer. The one-sided exact McNemar p-value in the registered direction is $p = 0.6128$. **Six tasks moved each way, so telling the model the exact number of hidden stages it had not seen did not reduce how often it claimed the task was finished.** Regenerated from `declare/data/` by `tables/make_tables.py`.

Reading the table: 115 tasks were played under both arms, and each task gives one of four answers. On 63 the model claimed at the decision point under both arms, and on 40 it claimed under neither, so those 103 tasks behave the same with the sentence and without it and carry no information about a difference. The information sits in the 12 tasks that changed answer, and they split 6 that claimed only under the control against 6 that claimed only under the warning, which is exactly the split chance produces when a sentence does nothing.

Statistical test:

Definition: Both arms played the same tasks, so each task has a pair of outcomes and only the tasks that disagree carry information about the arm. The null is that the warning has no effect, in which case each disagreement is equally likely to fall either way, and the probability of a split at least as extreme as the one observed comes from the exact binomial. This is McNemar’s test, one-sided in the registered direction.

Calculation: The split is 6 and 6 on 12 discordant pairs, and the one-sided exact value in the registered direction is $p = 0.6128$.

Inference: B1 is disconfirmed. The design had more than twice the discordant pairs it needed to fire, so this is not a test that failed for want of data. It is the absence of an effect.

Scoring the predictions: B1 disconfirmed, as above. B2 confirmed, since the warned model claimed 69 times at the decision point and was wrong 19 of them, where the registered clause needed one: knowing the exact count of hidden stages does not keep the claims that survive from being wrong. B3 confirmed, since the share of claims that were wrong tracked how often a green board means finished in each batch of the corpus, the split Table 7 prints arm by arm. B4 confirmed on both of its metrics, which read the window before the first decision point: the first is identical to four decimal places at 1.0261 submissions to the first green board under each arm, and the second is inside its band by 0.0016, which is one or two guesses out of sixty and is reported as a near miss rather than a comfortable pass. The warning therefore changed nothing about how the model worked, and nothing measurable about when it stopped.

The eliminated explanation: The information deficit was real, it was closed, and the behaviour did not move. A model that is told a task has K stages, that has seen fewer than K , and that could submit again for free, still says DONE. Whatever the failure is, it is not ignorance of how much is hidden.

The discrimination split: The corpus contains two kinds of task, and the model is never told which kind it is playing. A task that *hides a convention* is one whose hidden stages enforce a choice the description never states, so claiming at the decision point there is dangerous, while a task that *hides nothing* is one whose hidden stages add no new requirement, so passing everything visible really does mean finished and claiming there is safe. A secondary analysis, registered in the expectation of no difference, split every visit to the decision point by that kind, and it returned the largest effect in the programme: on the tasks where claiming is safe the model claimed at 97 of 117 visits, while on the tasks where claiming is dangerous it claimed at only 23 of 121. The difference is $+0.6390$ with a two-sided Fisher exact $p = 2.75e-24$, it replicates in both arms separately, and Table 4B prints the split.

The corpus-leak check: The corpus guarantees that the two kinds of task are indistinguishable in the bytes a solver sees on its first turn, and a leave-one-out classifier over those bytes confirms it, scoring 0.5185 against a base rate of 0.5259, which is a lift of -0.0074 . But a visit to the decision point happens later than turn 1. By then a stage has been released and the model has read that stage's tests, and on a task that hides a convention those tests encode the drawn convention while on a task that hides nothing they are quiet tests any visible-green solution already satisfies. The model is not seeing through the corpus. It is reading the tests it was just handed and judging whether its current code satisfies them, which is the job.

Inference: Two facts now sit together. The model can tell, at the decision point, whether claiming is safe, and it withholds the claim on roughly four fifths of the tasks that withhold something. On the remaining fifth it claims anyway. So the failure is in the threshold the model applies to a cue it reads correctly, rather than in the reading, and adding information to the rules did not move that threshold. Experiment 3 tries to move it with a payoff instead.

what the model was told it could do	kind of task	visits to the decision point	of those, it said DONE	share where it said DONE
may say DONE (control) (CTRL)	hides nothing	58	50	0.8621
may say DONE (control) (CTRL)	hides a convention	62	11	0.1774
told how many stages are hidden (WARN)	hides nothing	59	47	0.7966
told how many stages are hidden (WARN)	hides a convention	59	12	0.2034
both sets of rules pooled	hides nothing	117	97	0.8291
both sets of rules pooled	hides a convention	121	23	0.1901

Table 4B: **Experiment 2: the model can tell which tasks are safe to claim.** A task **hides a convention** when the code has to make a choice the task never states, several answers look reasonable, and the hidden tests accept exactly one; on such a task, passing everything visible does not mean the work is finished. A task that **hides nothing** is one where passing everything visible really does mean it is finished. Which kind a task is cannot be read from the bytes the model sees on its first turn: a leave-one-out classifier over those bytes scores 0.5185 against a base rate of 0.5259, a lift of -0.0074. A **visit to the decision point** is one moment where the model had passed everything visible with stages still hidden; a task can supply more than one. The 24-task annex is excluded because it carries no reliability label. The pooled difference is +0.6390, two-sided Fisher exact $p = 2.75e-24$. **The model claims far more often on tasks where claiming is correct, so it can tell at the decision point whether stopping is safe, and the failure is the threshold it applies to what it can tell rather than an inability to tell.** Regenerated from `declare/data/` by `tables/make_tables.py`.

7 Experiment 3: the incentive explanation

Bridge from Experiment 2: Experiment 2 ended with the model able to tell which tasks are risky while still claiming on a fifth of the risky ones, so the next candidate explanation is the criterion it applies. On that reading the model weighs the green board correctly but sets its threshold by the stakes, claiming when being wrong looks cheap, and a threshold set by the stakes must move when the stakes move. Experiment 3 therefore states a new price: it tells the model that stopping early is rewarded, and watches whether claiming rises.

Changes from Experiment 2: Less than anywhere else in the programme. The corpus, the frozen order and the model are unchanged, one new arm is played over the first 72 ids of that order, and the comparison is made against the CTRL record from Experiment 2 on the same ids. Two consequences follow: the comparison is again task by task between two arms that differ in one sentence, and the two runs come from different days, because the control was played in Experiment 2’s session, which is a weakness this design accepts and Experiment 4’s churn floor later calibrates.

The arm: SIGMA is CTRL plus one sentence appended to the rules, stating that a DONE

which turns out to be correct scores higher the fewer hidden stages the model has seen, and that a DONE which turns out to be wrong scores the same as running out of submissions. The sentence carries two clauses rather than one, because a reward with no stated cost is an instruction to claim rather than a payoff, and the second clause is true of this harness, since a wrong claim and an exhausted submission bound are both terminal and both unsolved. 72 episodes came back, at \$6.17.

The standing caveat: The payoff is stated and never delivered. Nothing in the harness pays a bonus and no score is ever shown to the model, so an effect here would be evidence that behaviour is sensitive to a described payoff structure rather than evidence about a real one. This caveat travels with every result in this section.

The registered power: Power is the probability that a design detects a shift of a given size, and the registration recorded this design's power in advance: 0.80 at a shift of 0.155, 0.50 at 0.10 and 0.18 at 0.05. In words, a ten-point shift in the claim rate was a coin flip for this design and a five-point shift was invisible, and the registration said in advance that both would land in the disconfirmed cell.

Predictions: Three were registered. C1: the stated payoff makes the model claim on more tasks, which is the incentive explanation put to the test, because a threshold set by the stakes must move when the stakes move. C2: the extra claims come from tasks where the model had been right to withhold, so accuracy among claims collapses, which is what a lowered threshold costs. C5: the working phase has not drifted between the two days on which the arms ran, tested, like Experiment 2's B4, on the window before the first decision point.

Results by arm: Under SIGMA the model played 72 tasks, said DONE at the decision point on 42 of them, and 10 of those claims were wrong, shares close to the control's on the same tasks, so the totals move nowhere. The registered view is again the paired one, task by task. Table 5 prints both, and the second panel of Figure 5 draws the tasks.

Reading the table: 71 tasks were played under both arms, and the four possible answers sort as they did in Experiment 2: on 41 the model claimed at the decision point under both arms, on 24 under neither, and those 65 tasks carry no information about the payoff. Six tasks changed answer, and five of the six changed against the prediction, by claiming under the control and not under the payoff, while one changed with it.

Statistical test:

Definition: This is the same paired test as Experiment 2, reading only the tasks whose answer differs between the two arms, and it is one-sided in the registered direction, which here is an increase in claiming under the payoff.

Calculation: The split is 5 against 1 on 6 discordant pairs, and the one-sided exact value in the registered direction is $p = 0.9844$.

Inference: C1 is disconfirmed. The p-value is as far on the wrong side of 0.5 as six observations can put it, and claiming fell rather than rose.

what the model was told it could do	tasks played	tasks where it said DONE at the decision point	of those, how many were wrong	share of tasks where it said DONE	tasks it claimed and the other did not
may say DONE (control) (CTRL)	118	71	21	0.6017	5
told that stopping early is rewarded (SIGMA)	72	42	10	0.5833	1

Table 5: **Experiment 3: telling the model that stopping early is rewarded does not make it claim more.** Both sets of rules were played over the same tasks, so the comparison is task by task rather than total against total. **71 tasks were played under both.** On 41 the model said DONE at the decision point under both sets of rules and on 24 it said DONE under neither; those 65 tasks carry no information about a difference and the registered test ignores them. That leaves 6 tasks that changed answer, split 5 against 1. The **share of tasks where it said DONE** column is the count of tasks with at least one DONE at the decision point divided by the tasks played in the column to its left; it is a description, and the registered test reads only the tasks that changed answer. The one-sided exact McNemar p-value in the registered direction is **p = 0.9844**. The two runs are from different days, which the churn-floor row of Table 6 exists to calibrate. The payoff was stated and never delivered: no bonus is paid and no score is shown to the model. This design was registered as able to detect a shift of 0.155 four times in five and a shift of 0.05 roughly one time in five, so what it can establish is a bound in one direction. **The rules stated that a correct early stop scored higher and that a wrong one cost no more than running out of submissions, and the model did not claim more often, which rules out a large increase in claiming from a stated reward and establishes nothing else.** Regenerated from declare/data/ by tables/make_tables.py.

The scope of the null: The reverse-direction test on the same six pairs returns $p = 0.1094$. That figure was not registered and is quoted only so it cannot later be told as though it had been, and it is quoted at all because it is the number that stops the decrease being over-read. Six discordant pairs cannot establish a direction either way. What this experiment licenses is a bound in one direction: a large increase in claiming is ruled out for the strongest payoff statement this harness can make while keeping the statement true, on this model, this corpus and this prefix, and nothing else is established.

The unclaimed headroom: 25 of the 71 tasks were not claimed at the decision point under the control, and the registration recorded that count in advance as the room an increase had to grow into. Under the payoff the model claimed on exactly one of them.

Scoring the predictions: C1 disconfirmed, as above. C2 disconfirmed, since accuracy among claims did not collapse but stayed level, at 0.2381 wrong per decision-point claim under the payoff against 0.2826 under the control. C5 confirmed on both of its

metrics, the second by one guess, so the two days did not drift in how the model worked and the paired comparison stands on that check.

Inference: Two descriptions of the payoff structure have now been given to this model, the stage count in Experiment 2 and the reward here, and neither moved the claim rate by an amount either design could see. Both put the information or the incentive in the rules, at the start of the episode, and asked the model to carry it to the decision point. Experiment 4 stops asking it to carry anything.

8 Experiment 4: the arithmetic at the decision point

Bridge from Experiment 3: The two failed interventions share a shape: each put a fact in the rules at the start of the episode, and each therefore required the model to retrieve that fact, hold it across the episode, and apply it many turns later at the moment it mattered. Neither tested whether the model performs that retrieval, so one explanation is still open, which is that the model has everything it needs and fails to bring it together at the moment of the decision. Experiment 4 removes the carrying requirement by putting the fact where the decision is made, already computed, and watches whether claiming falls.

Changes from Experiment 3: The corpus, the frozen order, the model and the 72 ids are all unchanged, and two new runs are played instead of one, both in the same session. That repairs the weakness Experiment 3 accepted, because the control is now played on the same day as its treatment, and it buys a second measurement for free: the control contract has now been played twice on the same tasks with nothing changed, so the pair of control runs measures how much this setup moves on its own, a quantity this paper calls the churn floor. The treatment also moves channel, from the rules the model reads once at the start to the feedback it reads at every turn, and that relocation is the entire intervention.

The two arms: CTRL' is the CTRL contract run again, changing nothing, so its job is twofold: it is the same-day control for the new arm, and against Experiment 2's CTRL run it supplies the churn floor. COUNT is CTRL' plus one line appended to the feedback at every release, at a fixed registered position: Released so far: stage R of K. Stages still withheld: N. R is the release index, K the task's stage count, and N the difference, so the line states, at the moment a fresh stage arrives, exactly how much remains hidden. This is the first arm in the programme whose treatment lives in the feedback channel rather than the rules. 72 episodes came back under CTRL' and 71 under COUNT, at \$12.51.

The information content of the line: It carries no information the warned model of Experiment 2 lacked, because K was in that arm's rules and R is in the model's own transcript. What it adds is the subtraction, already performed, in the channel the model is reading, at the moment the release creates the decision. If claiming falls here after not falling in Experiment 2, the difference between the two arms' fates cannot be the information, which was the same, and must be something about where and when it arrived.

Predictions: Three were registered. C4: the model claims at the decision point on fewer tasks under the counter, which is the intervention this section exists to test. C6: the churn floor comes in at or below the movement Experiment 2 recorded, meaning two identical runs disagree on no more of their tasks than the twelve of 115 that separated CTRL from WARN. C8: the counter changes the ending rather than the working, read on the same early window as B4 and C5 before it.

Results by arm: Under CTRL' the model played 72 tasks, said DONE at the decision point on 44 of them, and 12 of those claims were wrong. Under COUNT it played 71 tasks, claimed on 31, and 6 were wrong, so for the first time in the programme the totals move: the share of tasks claimed falls from 0.6111 to 0.4366, and the wrong claims halve. Table 6 prints the two arms with the churn floor as a third row, the third panel of Figure 5 draws the tasks, and its fourth panel draws the churn floor on the same scale.

Reading the table: 71 tasks were played under both arms, and the four answers sort one more time: on 30 the model claimed at the decision point under both arms, on 27 under neither, and those 57 tasks carry no information about the counter. Fourteen tasks changed answer, and thirteen of the fourteen changed by ceasing to claim under the counter, while one changed the other way.

Statistical test:

Definition: The same paired form again, one-sided in the registered direction, which here is a decrease.

Calculation: The split is 13 against 1 on 14 discordant pairs, and the one-sided exact value in the registered direction is $p = 0.0009$.

Inference: C4 is confirmed. The realised shift in per-task claiming is 0.1745 against a design registered as powered for 0.155, so this is the first effect in the programme that its design was powered to see. Wrong claims at the decision point fell from 12 to 6.

The churn floor: C6 bought the number the programme had been missing. The control contract was run a second time on the same ids, changing nothing at all, and the two identical runs disagreed with each other on 2 tasks of 71. That is how much this setup moves on its own. Against it, Experiment 2's twelve of 115 and Experiment 3's six of 71 are three to four times the noise, and this experiment's fourteen is seven times it. The floor does not rescue the two nulls, since a movement above the noise floor that splits evenly is still an even split. What it does is establish that those counts were large enough to have shown a direction had one existed.

The composition of the suppressed claims: Of the thirteen tasks the counter stopped the model claiming on, 8 were tasks that hide nothing, where the claim would have been true, and 5 were tasks that hide a convention, where it would have been false. The single task that moved the other way hides nothing and its claim was true, and it is also one of the two tasks on which the two identical control runs disagreed, so the control is the odd run out on it rather than the counter. That is two events and no test, and it is recorded because a thirteen-against-one split reads differently if its one

what the model was told it could do	tasks played	tasks where it said DONE at the decision point	of those, how many were wrong	share of tasks where it said DONE	tasks it claimed and the other did not
may say DONE (control, re-run) (CTRL')	72	44	12	0.6111	13
shown the count of what is still hidden (COUNT)	71	31	6	0.4366	1
the same control played a second time, against its first playing (CTRL' vs CTRL)	71	-	-	-	2

Table 6: **Experiment 4: doing the arithmetic at the decision point does change what the model does.** Both sets of rules were played over the same tasks, so the comparison is task by task rather than total against total. **71 tasks were played under both.** On 30 the model said DONE at the decision point under both sets of rules and on 27 it said DONE under neither; those 57 tasks carry no information about a difference and the registered test ignores them. That leaves 14 tasks that changed answer, split 13 against 1. The **share of tasks where it said DONE** column is the count of tasks with at least one DONE at the decision point divided by the tasks played in the column to its left; it is a description, and the registered test reads only the tasks that changed answer. The one-sided exact McNemar p-value in the registered direction is **p = 0.0009**. The final row is the churn floor: the same control contract played a second time on the same tasks, changing nothing, disagreed with itself on 2 of 71 tasks, so 2 is how much this setup moves on its own. The line reads Released so far: stage R of K. Stages still withheld: N. and carries no information the model in Experiment 2 was not already in a position to work out; what it adds is the subtraction, already done, in the channel the model is reading. This design cannot separate the content of the line from the presence of an extra sentence, and the placebo arm that would separate them has not been run. **One line of feedback, printed at the moment the model decides, cut how often the model claimed the task was finished and halved the wrong claims.** Regenerated from declare/data/ by tables/make_tables.py.

counter-example is a coin flip. Table 6B prints the accounting.

Conservatism and accuracy: The counter did not make the model better at guessing the withheld convention. It made it stop more often, and stopping more often is disproportionately rewarded on this harness because submitting again is free. The model solved more tasks rather than fewer, since correct claims went from 60 to 63 and tasks solved by the bound from 60 to 64, on one fewer episode, and Table 6B prints both movements. Giving up eight correct early stops cost nothing in the task's own terms, and that is a property of this harness rather than a general result, since a harness that charged for submissions would price the trade differently.

The unrun placebo control: This design cannot separate the content of the counter

quantity	count
tasks the counter stopped the model claiming on	13
of those, tasks that hide nothing, where the claim would have been true	8
of those, tasks that hide a convention, where the claim would have been false	5
tasks that moved the other way (L5-099, which hides nothing, and is one of the two churn-floor tasks)	1
correct DONEs, control against counted	60 → 63
tasks solved by the submission bound, control against counted	60 → 64

Table 6B: **Experiment 4: the destination of the flipped tasks and the price of the counter.** The thirteen tasks the counter stopped the model claiming on split 8 that hide nothing, where the suppressed claim would have been true, against 5 that hide a convention, where it would have been false, so the counter suppressed good and bad claims together rather than sharpening the model’s guess about the withheld convention. The one task that moved the other way is one of the two tasks on which the two identical control runs disagreed, so the control is the odd run out on it rather than the counter. The last two rows price the trade: correct claims and solved tasks both rose, because a suppressed correct claim converts into a solve by the submission bound at no cost on this harness. **The counter bought conservatism rather than accuracy, and the purchase was free only because submitting again costs nothing here.** Regenerated from `declare/data/` by `tables/make_tables.py`.

line from the presence of an extra sentence of feedback. Both would produce this result. The mechanism this experiment is meant to license is that the failure is one of integration rather than of threshold, and that mechanism follows only if the content of the line is doing the work. A placebo arm carrying a line of the same shape with its numbers removed is the missing control. It was not bought.

Inference: The same information, moved from the rules to the feedback and from the start of the episode to the moment of the decision, changed the behaviour that two earlier arms could not move. Read alongside Experiment 2, where the model already discriminated the risky tasks from the safe ones, the result suggests the model has the parts it needs and does not combine them at the moment it decides unless the combination is supplied. That reading rests on one experiment whose control has not been run.

9 The four experiments together

Questions and answers: Figure 6 draws the programme as a tree: the four experiments in time order, the fourteen registered predictions under them, and at the bottom the ledger of explanations each experiment removed. The questions and answers follow in the order the programme ran them.

- Experiment 1 asked whether a model given the claim channel uses it on work it

cannot verify. It does, and it does so more often as the model gets stronger.

- Experiment 2 asked whether the cause is missing information. It is not, because supplying the exact count of hidden stages changed nothing.
- Experiment 3 asked whether the cause is the incentive. Within the bound that design could see, it is not.
- Experiment 4 asked whether supplying the arithmetic at the moment of the decision changes the behaviour. It does, and it is the only intervention in the programme that did.




Figure 6: How the programme unravelled. The top row is the four experiments in time order, each naming its arms with a one-line description of what each arm was told. The middle row is the fourteen registered predictions, marked with the verdict their registered analysis returned. The bottom row is the explanation ledger: what each experiment eliminated, and the one explanation left standing, whose placebo control has not been run. Drawn by `figures/make_figures.py:f3_timeline` from `tables/T8.csv` and `declare/data/runs.csv`.

The location of the failure: Across 856 episodes and 608 claims, the programme produced 74 wrong claims, of which 73 were made at the decision point, while claims made once every stage had been released were never wrong, 344 times out of 344, and no claim was made after a submission that left a visible test failing, across 115 opportunities. The error is therefore concentrated at the one state where the visible evidence is complete and the hidden evidence is not.

The reliability of the green board: Table 7 splits the wrong claims by how much a green board is worth, and that worth is a property of the corpus: it was built in two batches, one in which a quarter of the tasks hide nothing and one in which three quarters do, so a board with every visible test passing is usually a lie in the first batch

and usually the truth in the second. The model is never told which batch a task came from, and the batch cannot be read from the bytes it sees. Claims turn out to be wrong far more often in the batch where a green board is less often trustworthy, which means the model is weighting a real cue by a reliability it has no way to observe, and nothing available to it inside an episode would let it calibrate that weight.

experiment	what the model was told it could do	how often a green board means finished	DONEs at the decision point	of those, how many were wrong	share of its DONEs that were wrong
experiment 2	may say DONE (control) (CTRL)	1 task in 4	19	9	0.4737
experiment 2	may say DONE (control) (CTRL)	3 tasks in 4	42	2	0.0476
experiment 2	told how many stages are hidden (WARN)	1 task in 4	18	7	0.3889
experiment 2	told how many stages are hidden (WARN)	3 tasks in 4	41	2	0.0488
experiment 3	told that stopping early is rewarded (SIGMA)	1 task in 4	9	3	0.3333
experiment 3	told that stopping early is rewarded (SIGMA)	3 tasks in 4	27	1	0.0370
experiment 4	may say DONE (control, re-run) (CTRL')	1 task in 4	13	5	0.3846
experiment 4	may say DONE (control, re-run) (CTRL')	3 tasks in 4	25	1	0.0400
experiment 4	shown the count of what is still hidden (COUNT)	1 task in 4	6	1	0.1667
experiment 4	shown the count of what is still hidden (COUNT)	3 tasks in 4	20	0	0.0000
all four experiments pooled		1 task in 4	65	25	0.3846
all four experiments pooled		3 tasks in 4	155	6	0.0387

Table 7: **How often a DONE was wrong, against how much a green board is worth.** The corpus was built in two batches that differ in how often a task hides nothing. In the first batch **1 task in 4 hides nothing**, so a board with every visible test passing is usually a lie; in the second **3 tasks in 4 hide nothing**, so the same board is usually the truth. The model is never told which batch a task came from, and the batch cannot be read from the bytes it sees. Each row counts only the DONEs said at the decision point, on core tasks; the 24-task annex carries no batch label and is excluded. The pooled difference between the two batches is 25/65 against 6/155, two-sided Fisher exact $p = 2.22e-10$. **A DONE is wrong far more often in the batch where a green board is less often trustworthy, which means the model's claim is tracking a cue whose reliability it cannot observe.** Regenerated from `declare/data/` by `tables/make_tables.py`.

Eliminated and surviving explanations: Table 8 lists all 14 registered predictions with their verdicts: 8 confirmed, 5 refuted and 1 returning neither, and the refutations are what carry the argument. Missing information was eliminated, because supplying the exact stage count moved nothing; incentive was eliminated within the bound that design could see; and perception was eliminated, because the model already discriminates the two kinds of task by a margin of +0.6390. What survives is integration, meaning the step in which a model combines what its rules told it with what its own transcript shows, at the moment it decides.

experiment	label	what was registered before the run	verdict	what the data showed
experiment 1	P1	a model given the channel claims completion it cannot verify, at the hardest difficulty	CONFIRMED	4 wrong claims by the strongest model, all in the hardest band
experiment 1	P2	the benefit of removing the channel is largest for the middle model	NOT CONFIRMED	the ordering holds, the test rests on one task that changed answer
experiment 1	P3	removing the channel helps less where tests are staged	DISCONFIRMED	it helped more: 7 tasks rescued of 96 against 1 of 107
experiment 1	P4	the benefit tracks the count of wrong claims	DISCONFIRMED	the two orderings disagree across three models
experiment 2	B1	telling the model how many stages are hidden reduces claiming	DISCONFIRMED	6 tasks moved each way of 115, $p = 0.6128$
experiment 2	B2	wrong claims survive an explicit stage count	CONFIRMED	69 claims under the warning, 19 of them wrong
experiment 2	B3	accuracy tracks how much a green board is worth	CONFIRMED	0.4324 wrong against 0.0482, $p < 0.0001$
experiment 2	B4	the warning changes the ending, not the working	CONFIRMED	both checks inside their bands, one by 0.0016
experiment 3	C1	a stated reward for stopping early makes the model claim more	DISCONFIRMED	6 tasks moved, 5 against the prediction, $p = 0.9844$
experiment 3	C2	the extra claims come from where it was right to withhold	DISCONFIRMED	accuracy did not collapse: 0.2381 against 0.2826
experiment 3	C5	the working phase has not drifted between the two days	CONFIRMED	both checks inside their bands
experiment 4	C4	doing the arithmetic at the decision point reduces claiming	CONFIRMED	14 tasks moved, 13 against 1, $p = 0.0009$
experiment 4	C6	two identical runs disagree no more than experiment 2 did	CONFIRMED	2 tasks of 71 = 0.0282 against 0.10435
experiment 4	C8	the counter changes the ending, not the working	CONFIRMED	+0.0000 and +0.0278, both inside band

Table 8: **Every registered prediction and how it turned out.** Every prediction the programme registered, with the verdict its own registered analysis returned. A prediction was written down, together with the result that would confirm it and the result that would refute it, and frozen in a version-controlled file before the first episode of its experiment ran. **CONFIRMED** means the registered confirmation clause was met, **DISCONFIRMED** means the registered refutation clause was met, and **NOT CONFIRMED** means neither was met, which is possible because one early registration’s two clauses did not cover the whole range of outcomes. The verdicts are transcribed from the four write-ups and are not recomputed here, because a verdict is a reading of a clause rather than a number. **Of 14 registered predictions 8 were confirmed, 5 were refuted and 1 returned neither, and the refutations are what removed the explanations this paper eliminates.** Regenerated from `declare/data/by tables/make_tables.py`.

The proposed mechanism: The model reads a green board as evidence of completion and applies a threshold to it. That threshold is too low for the tasks that hide something, it is not moved by putting information about the hidden stages in the rules, and it is not moved by attaching a payoff to stopping, yet the same information moves it when it arrives already computed, in the channel the model is reading, at the moment of the decision. The claim this paper makes is that the failure is one of integration at the decision point, and the evidence for that claim is one experiment whose placebo control has not been run.

The scope of Table 9: The five sets of rules that played the same 72 tasks are listed as rows for reference. Three of them ran on one day and two on another, and a difference between rows from different days cannot be told apart from a difference between the days, so no comparison is drawn from that table. Every comparison in this paper is one of the task-by-task ones in Tables 4, 5 and 6, each reading only tasks played under both of the arms it compares.

experiment	what the model was told it could do	day	tasks played	visits to the decision point	DONEs at the decision point	of those, how many were wrong	share of tasks where it said DONE
experiment 2	may say DONE (control) (CTRL)	2026-08-09	71	87	46	13	0.6479
experiment 2	told how many stages are hidden (WARN)	2026-08-09	72	91	45	11	0.6250
experiment 3	told that stopping early is rewarded (SIGMA)	2026-08-10	72	91	42	10	0.5833
experiment 4	may say DONE (control, re-run) (CTRL)	2026-08-10	72	88	44	12	0.6111
experiment 4	shown the count of what is still hidden (COUNT)	2026-08-10	71	96	31	6	0.4366

Table 9: **The five sets of rules on the same 72 tasks.** All five sets of rules restricted to the same 72 tasks, so the rows describe the same work. **These rows are a description and not a comparison.** Three of them were played on one day and two on another, and a difference between two rows from different days cannot be told apart from a difference between the two days. The comparisons this paper draws are the task-by-task ones in Tables 4, 5 and 6, each of which reads only tasks played under both sets of rules being compared. **The table is included so that the five sets of rules can be seen side by side, and no inference is drawn from any difference between its rows.** Regenerated from `declare/data/by tables/make_tables.py`.

10 Limitations

Every limitation below is taken from the programme’s own registrations and deviation logs rather than assembled for this section, and each one names what it does and does not put at risk.

1. **The placebo arm has not been run:** Experiment 4 cannot separate the content of the counter line from the presence of an extra sentence of feedback. The registration recorded this before the run and the secondary analysis written to bear on it does not settle it. It is the largest open question in the paper, and the mechanism proposed in Section 9 depends on it.
2. **One model carries three of the four experiments:** Experiments 2, 3 and 4 are all the same frontier model, so every elimination rests on one model. The capability result in Section 5 is three models, and it is the only claim here that is not single-model.

3. **One provider and one snapshot of it:** Nothing here is replicated across providers, and a different frontier model could in principle behave differently at the decision point.
4. **The decisive experiment rests on 71 paired tasks:** Effects in this programme live in counts of single digits, which is why every test is exact and why no confidence interval appears anywhere. A thirteen-against-one split is a strong result at this size and would still be a small sample at any other.
5. **Two of the comparisons cross a day boundary:** Experiment 3's arm ran the day after its control, and the churn floor calibrates this once, at 2 tasks of 71, but the floor says what a third identical run would probably do rather than what a different day would do, and it was measured only in the same session.
6. **The corpus is generated and its conventions come from a catalogue:** Whether the same failure appears on organic under-specification in real repositories is untested, and the guarantee that makes this instrument work is exactly the guarantee a real repository cannot give.
7. **Experiment 1's registration did not partition the outcome space:** One prediction is recorded as neither confirmed nor disconfirmed because its two clauses left a gap and the run landed in it. Later registrations fixed this by requiring three clauses that partition.
8. **Episodes that never came back are dropped rather than counted:** Four in Experiment 2 and one each elsewhere failed in transport, and each is excluded from the paired statistics rather than scored as a failure, because an absent episode is not an unsolved one.
9. **Operational deviations are logged rather than repaired:** One publish gate was waived and recorded verbatim as a deviation, and a hash taken over a directory listing rather than over a named file set meant that two pooled figures in Experiment 2 print differently once later result files exist. The published values are quoted here, the population they are taken over is named, and no verdict moves either way.

11 Design implications and reproducibility

Design advice: The advice this programme licenses is narrower than the previous paper's, and it should be stated as such. If an agent must judge for itself when a task is finished, put the state it needs in front of it at the moment it decides, already computed. Putting the same fact in the system prompt did not work across 115 paired tasks, attaching a payoff to stopping early did not work within a bound that rules out a large increase, and printing one line of arithmetic into the feedback at the moment of the decision moved 13 tasks against 1 and halved the wrong claims.

Two qualifications travel with that. The first is that the placebo arm has not run,

so a designer following this advice is following a result whose mechanism is not established. The second is that the intervention bought conservatism rather than accuracy, and it was free here only because submitting again costs nothing on this harness. A system that charges for a retry, whether in latency, tokens or a user’s patience, will price that trade differently and should measure it before adopting the line.

The quantity to instrument: The failure has a location and the location is cheap to detect, because it is the state where every check the agent can run passes and the agent knows more checks exist. A harness that logs nothing else can log that state and count what its model does there, and that count is the quantity everything in this paper turns on.

Scale and cost: The four experiments comprise 856 episodes over three models, the eleven runs cost \$55.21, and the whole programme including calibration probing that produced no episode cost \$61.25. An experiment that locates a reliability failure and then tries three interventions against it fits inside a rounding error on a compute budget, which is worth knowing before choosing to ship a harness on intuition instead.

Reproducibility: Every number in this paper is recomputed from committed episode logs by one command, reading only frozen result files, task metadata and the checker, with no model call and no key. The whole programme is also published as a flattened dataset of five tables and 9,194 rows, extracted once and deterministically from the seventeen committed result files, so a reader can ask a question with a `GROUP BY` rather than by rerunning anything. Both corpora are pinned by content hash, each experiment’s predictions are frozen at a commit that provably precedes its first live call, and every deviation is recorded in an append-only log.

Provenance: The harness, scaffolding and analysis code were built with AI assistance under execution-verified acceptance gates, with every phase advancing only on raw command output audited by the author and all deviations logged. The research questions, the registered predictions, the verification discipline and the interpretations are the author’s.

12 References

- McNemar, Q. (1947). Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika*, 12(2), 153–157.
- Fisher, R. A. (1935). The logic of inductive inference. *Journal of the Royal Statistical Society*, 98(1), 39–82.
- Bhatnagar, K. (2026). Who decides when the task is done? Measuring the effect of moving completion authority from LLMs to checkers. DOI 10.5281/zenodo.21698264.